

A Study of Optimal 4-Bit Reversible Toffoli Circuits and Their Synthesis

Oleg Golubitsky and Dmitri Maslov, *Member, IEEE*

Abstract—Optimal synthesis of reversible functions is a nontrivial problem. One of the major limiting factors in computing such circuits is the sheer number of reversible functions. Even restricting synthesis to 4-bit reversible functions results in a huge search space ($16! \approx 2^{44}$ functions). The output of such a search alone, counting only the space required to list Toffoli gates for every function, would require over 100 terabytes of storage. In this paper, we present two algorithms: one, that synthesizes an optimal circuit for any 4-bit reversible specification, and another that synthesizes all optimal implementations. We employ several techniques to make the problem tractable. We report results from several experiments, including synthesis of all optimal 4-bit permutations, synthesis of random 4-bit permutations, optimal synthesis of all 4-bit linear reversible circuits, and synthesis of existing benchmark functions; we compose a list of the hardest permutations to synthesize, and show distribution of optimal circuits. We further illustrate that our proposed approach may be extended to accommodate physical constraints via reporting LNN-optimal reversible circuits. Our results have important implications in the design and optimization of reversible and quantum circuits, testing circuit synthesis heuristics, and performing experiments in the area of quantum information processing.

Index Terms—Reversible circuits, logic synthesis, quantum circuits.

1 INTRODUCTION

To the best of our knowledge, at present, physically reversible technologies are found only in the quantum domain [12]. However, “quantum” unites several technological approaches to information processing, including ion traps, optics, superconducting, spin-based and cavity-based technologies [12]. Of those, trapped ions [5] and liquid state Nuclear Magnetic Resonance (NMR) [13] are two of the most developed quantum technologies targeted for computation in the circuit model (as opposed to communication or adiabatic computing). These technologies allow computations over a set of 8 qubits and 12 qubits, correspondingly.

Reversible circuits are an important class of computations that need to be performed efficiently for the purpose of efficient quantum computation. Multiple quantum algorithms contain arithmetic units such as adders, multipliers, exponentiation, comparators, quantum register shifts, and permutations, that are best viewed as reversible circuits. Moreover, reversible circuits are indispensable in quantum error correction [12]. Often, the efficiency of the reversible implementation is the bottleneck of a quantum algorithm (e.g., integer factoring and discrete logarithm [19]) or even a class of quantum circuits (e.g., stabilizer circuits [1]).

In this paper, we report algorithms that find optimal circuit implementations for 4-bit reversible functions. These algorithms have a number of potential uses and implications.

One major implication of this work is that it will help physicists with experimental design, since foreknowledge of the optimal circuit implementation aids in the control over quantum mechanical systems. The control of quantum mechanical systems is very difficult, and as a result experimentalists are always looking for the best possible implementation. Having an optimal implementation helps to improve experiments or show that more control over a physical system needs to be established before a certain experiment could be performed. To use our results in practice requires defining minimization criteria (e.g., implementation cost of gates versus depth versus architecture, etc.) dictated by a particular technology used, that may differ from one quantum information processing approach to another. Consequently, in this paper, we ignored such physical constraints, but concentrated on the minimization of the gate count. This serves as a proof of principle, showing that the search is possible in any practical scenario. We further explain how to modify our algorithms to account for more complex circuit minimization criteria in Section 6, and illustrate one of such modifications in the Section 7.

A second important contribution is due to the efficiency of our implementation—0.00756 seconds per synthesis of an optimal 4-bit reversible circuit. The algorithm could easily be integrated as part of peephole optimization, such as the one presented in [16].

Furthermore, our implementation allows to develop a subset of optimal implementations that may be used to test heuristic synthesis algorithms. Currently, similar tests are performed by comparison to optimal 3-bit implementations [4], [6], [8]. The best heuristic solutions have very tiny overhead when compared to optimal implementations,

• O. Golubitsky is with Google, Inc., 427A Dunvegan Dr., Waterloo, ON N2K 2C8, Canada. E-mail: oleg.golubitsky@gmail.com.

• D. Maslov is with the US National Science Foundation, Institute for Quantum Computing and the Department of Physics and Astronomy, University of Waterloo, 200 University Ave. W, Waterloo, ON N2L 3G1, Canada. E-mail: dmitri.maslov@gmail.com.

Manuscript received 1 Feb. 2011; revised 15 June 2011; accepted 12 July 2011; published online 22 July 2011.

Recommended for acceptance by I. Markov.

For information on obtaining reprints of this article, please send e-mail to: tc@computer.org, and reference IEEECS Log Number TC-2011-02-0075.

Digital Object Identifier no. 10.1109/TC.2011.144.

making such a test hard to improve. As such, it would help to replace this test with a more difficult one that allows more room for improvement. We suggest that this test set should include known benchmarks, and a combination of other functions—linear reversible as well as, possibly, representatives from other classes, those with a few gates and those requiring a large number of gates, etc. We have not worked out the details of such a test.

Finally, due to the effectiveness of our approach, we are able to report new optimal implementations for small benchmark functions, calculate $L(4)$, the number of reversible gates required to implement a reversible 4-bit function, calculate the average number of gates required to implement a 4-bit permutation, and show the distribution of the number of permutations that may be implemented with the predefined number of gates.

An earlier version of this paper has been presented at the DAC'2010 conference.

2 PRELIMINARIES

2.1 Quantum Computing

We start with a very short review of basic concepts in quantum computing. An in-depth coverage may be found in [12].

The state of a single qubit is described by a linear combination (Dirac notation)/column vector $\alpha|0\rangle + \beta|1\rangle = (\alpha, \beta)^t$, where α and β are complex numbers called the amplitudes, and $|\alpha|^2 + |\beta|^2 = 1$. Real numbers $|\alpha|^2$ and $|\beta|^2$ represent the probabilities of reading the logic states $|0\rangle$ and $|1\rangle$ upon (computational basis) measurement. The state of a quantum system with n qubits is described by an element of the tensor product of the single state spaces and can be represented as a normalized vector of length 2^n , called the state vector. Furthermore, quantum mechanics allows evolution of the state vector through its multiplication by $2^n \times 2^n$ unitary matrices called the gates. These gates may be applied to a quantum state sequentially—such process constitutes constructing a circuit—which is equivalent to a series of proper matrix multiplications. To illustrate the gate application, take the two qubit state vector $|11\rangle = (0, 0, 0, 1)^t$ and apply a CNOT gate, defined as the matrix

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{bmatrix}.$$

The result is the state $|10\rangle = (0, 0, 1, 0)^t$. It may be observed that in the Dirac notation the CNOT gate may be described as follows: application of the CNOT gate flips the value of the second qubit iff the value of the first qubit is one. Dropping the bra-ket Dirac notation results in the following redefinition over Boolean values a and b —gate CNOT performs transformation $a, b \mapsto a, b \oplus a$. This definition extends to vectors by the linearity, and thus is not ambiguous. In particular, in the follow up sections, we will consider reversible circuits (sometimes known as quantum Boolean circuits)—those where the matrix entries are strictly Boolean/integer, and for simplicity we will drop bra and ket in the notations, leaving just the variable names.




Fig. 1. NOT, CNOT, Toffoli, and Toffoli-4 gates.

The set of reversible circuits forms a group, that is also a subgroup of the set of all unitary transformations.

2.2 Reversible Circuits

In this paper, we consider circuits with NOT, CNOT, Toffoli (TOF), and Toffoli-4 (TOF4) gates defined as follows:

- NOT(a) : $a \mapsto a \oplus 1$.
- CNOT(a, b) : $a, b \mapsto a, b \oplus a$.
- TOF(a, b, c) : $a, b, c \mapsto a, b, c \oplus ab$.
- TOF4(a, b, c, d) : $a, b, c, d \mapsto a, b, c, d \oplus abc$.

where \oplus denotes the EXOR operation and concatenation is the Boolean AND, see Fig. 1 for illustration. These gates are used widely in quantum circuit construction, and have been demonstrated experimentally in multiple quantum information processing proposals [12]. In particular, CNOT is a very popular gate among experimentalists, frequently used to demonstrate control over a multiple-qubit quantum mechanical system. Since quantum circuits describe time evolution of a quantum mechanical system where individual “wires” represent physical instances, and time propagates from left to right, this imposes restrictions on the circuit topology. In particular, quantum and reversible circuits are strings of gates. As a result, feedback (time wrap) is not allowed and there may be no fan-out (mass/energy conservation).

In this paper, we are concerned with searching for circuits requiring a minimal number of gates. Our focus is on the proof of principle, i.e., showing that any optimal 4-bit reversible function may be synthesized fast, rather than attempting to report optimal implementations for a number of potentially plausible cost metrics. In fact, our implementation allows other circuit cost metrics to be considered, as discussed in Sections 6 and 7.

In related work, there have been a few attempts to synthesize optimal reversible circuits with more than three inputs. Große et al. [3] employ SAT-based technique to synthesize provably optimal circuits for some small parameters. However, their implementation quickly runs out of resources. The longest optimal circuit they report contains 11 gates. The latter took 21,897.3 seconds to synthesize—same function that the implementation we report in this paper synthesized in 0.000052 seconds, see Table 7. Prasad et al. [16] used breadth first search to synthesize 26,000,000 optimal 4-bit reversible circuits with up to six gates in 152 seconds. We extend this search into finding all $16!$ optimal circuits in 1,130,276 seconds. This is over 100 times faster (per circuit) and 800,000 times more than reported in [16]. Yang et al. [20] considered short optimal reversible 4-bit circuits composed with NOT, CNOT, and Peres [14] gates. They were able to synthesize optimal circuits with up to six gates, and use those to optimally synthesize any given even permutation requiring no more than 12 gates. In other




Fig. 2. (a) A suboptimal. (b) An optimal circuit for 1-bit full adder.

words, they can search a space of the size equal to approximately one quarter of the number of all 4-bit reversible functions. Our algorithms and implementation allow optimal synthesis of *all* 4-bit reversible functions and *any* 4-bit reversible function, and it is much faster.

2.3 Motivating Example

Consider the two reversible circuit implementations in Fig. 2 of a 1-bit full adder. This elementary function/circuit serves as a building block for constructing integer adders. The famous Shor's integer factoring algorithm is dominated by adders like this. As such, the complexity of an elementary 1-bit adder circuit largely affects the efficiency of factoring an integer number with a quantum algorithm. It is, thus, important to have a well-optimized implementation of a 1-bit adder as well as other similar small quantum circuit building blocks.

In this paper, we consider the synthesis of optimal circuits, i.e., we provably find the best possible implementation. Using optimal implementations of circuits potentially increases the efficiency of quantum algorithms and helps to reduce the difficulty with controlling quantum experiments.

3 FINDOPT: AN ALGORITHM TO FIND AN OPTIMAL CIRCUIT

We first outline our algorithm for finding an optimal circuit and then discuss it in detail in the follow up sections.

There are $N = 2^n$ reversible n -variable functions. The most obvious approach to the synthesis of all optimal implementations is to compute all optimal circuits and store them for later look-up. However, this is extremely inefficient. This is because such an approach requires $\Omega(N)$ space and, as a result, at least $\Omega(N)$ time. These space and time estimates are lower bounds, because, for instance, storing an optimal circuit requires more than a constant number of bits, but for simplicity, let us assume these figures are exact. Despite considering both figures for space and time impractical, we use this simple idea as our starting point.

We first improve the space requirement by observing that if one synthesized all halves of all optimal circuits, then it is possible to search through this set to find both halves of any optimal circuit. It can be shown that the space requirement for storing halves has a lower bound of $\Omega(\sqrt{N})$. However, searching for two halves potentially requires a runtime on the order of the square of the search space, $\Omega((\sqrt{N})^2) = \Omega(N)$, a figure for runtime that we deemed inefficient. Our second improvement is, thus, to use a hash table to store the optimal halves. This reduces the runtime to soft $\Omega(\sqrt{N})$. While this lower bound does not necessarily imply that the actual complexity is lower than $O(N)$, this turns out to be

the case, because the set of optimal halves is indeed much smaller than the set of all optimal circuits (an analytic estimate for the relative size of the former set is hard to obtain, though). Cumulatively, these two improvements reduce $\Omega(N)$ space and $\Omega(N)$ time requirement to $O(\#halves(N))$ space and soft $O(\#halves(N))$ time requirement. These reductions almost suffice to make the search possible using modern computers.

Our last step, apart from careful coding, that made the search possible is the reduction of the space requirement (with consequent improvement for runtime) by a constant of almost 48 via exploiting the following two features. First, simultaneous input/output relabeling, of which there are at most $24(=4!)$ different ones, does not change the optimality of a circuit. And second, if an optimal circuit is found for a function f , an optimal circuit for the inverse function, f^{-1} , can be obtained by reversing the optimal circuit for f . This allows to additionally “pack” up to twice as many functions into one circuit. The cumulative improvement resulting from these two observations, is by a factor of almost $2 \times 24 = 48$. Due to symmetries, the actual number is slightly less. See Table 4 (Column 2 versus Column 3) for exact comparison.

3.1 The Search-and-Lookup Algorithm

For brevity, let the size of a reversible function mean the minimal number of gates required to implement it. Using breadth-first search, we can generate the smallest circuits for all reversible functions of size at most k , for a certain value of k . (This can be done in advance, on a larger machine, and need not be repeated for each reversible function.)

Assume that the given function f , for which we need to synthesize a minimal circuit, has size at most $2k$. We can first check whether f is among the known functions of size at most k and, if so, output the corresponding minimal circuit. If not, then the size of f is between $k+1$ and $2k$, inclusive, and there exist reversible functions h and g of size k and at most k , respectively, such that $f = h \circ g$. If we find such g of the smallest size, then we can obtain the smallest circuit for f by composing the circuits for g and h .

Multiplying the above equality by g^{-1} , we obtain $f \circ g^{-1} = h$. Observe that g^{-1} has the same size as g . Therefore, by trying all functions g of size $1, 2, \dots, k$ until we find one such that $f \circ g$ has size k , we can find a g of the smallest size.

The above algorithm involves sequential access to the functions of size at most k and their minimal circuits and a membership test among functions of size k . Since the latter test must be fast and requires random memory access, we need to store all functions of size k in the memory. Thus, the amount of available RAM imposes an upper bound on k .

In practice, we store a 4-bit reversible function using a 64-bit word, because this allows for an efficient implementation of functional composition, inversion, and other necessary operations. On a typical PC with 4 GB of RAM, we can store all functions for $k = 6$. This means that we can apply the above search algorithm only to functions of size at most 12. Unfortunately, this will not cover all 4-bit reversible functions. Therefore, further reduction of memory usage is necessary.

3.2 Symmetries

A significant reduction of the search space can be achieved by taking into account the following symmetries of circuits:

1. Simultaneous relabeling of inputs and outputs. Given an optimal circuit implementing a 4-bit reversible function f with inputs x_0, x_1, x_2, x_3 and outputs y_0, y_1, y_2, y_3 and a permutation $\sigma : \{0, 1, 2, 3\} \rightarrow \{0, 1, 2, 3\}$, we can construct a new circuit by relabeling the inputs and outputs into $x_{\sigma(0)}, x_{\sigma(1)}, x_{\sigma(2)}, x_{\sigma(3)}$ and $y_{\sigma(0)}, y_{\sigma(1)}, y_{\sigma(2)}, y_{\sigma(3)}$, respectively. Then, the new circuit will provide a minimal implementation of the corresponding reversible function f_σ . Indeed, if it is not minimal and there is an implementation of f_σ by a circuit with a smaller number of gates, we can relabel the inputs and outputs of this implementation with σ^{-1} and obtain a smaller circuit implementing the original function f . This contradicts the assumption that the original circuit for f is optimal.

Given f and σ , a formula for f_σ can be easily obtained. Observe that the mapping $x_0, x_1, x_2, x_3 \mapsto x_{\sigma(0)}, x_{\sigma(1)}, x_{\sigma(2)}, x_{\sigma(3)}$ is a 4-bit reversible function, which we denote by $\tilde{\sigma}$. The mapping $y_{\sigma(0)}, y_{\sigma(1)}, y_{\sigma(2)}, y_{\sigma(3)} \mapsto y_0, y_1, y_2, y_3$ is then given by the inverse, $\tilde{\sigma}^{-1}$. Therefore, the 4-bit values y_0, y_1, y_2, y_3 of f_σ on a 4-bit tuple x_0, x_1, x_2, x_3 can be obtained by applying first $\tilde{\sigma}$, then f , and finally $\tilde{\sigma}^{-1}$. We obtain $f_\sigma = \tilde{\sigma}^{-1} \circ f \circ \tilde{\sigma}$. We call the set of functions f_σ the *conjugacy class* of f modulo simultaneous input/output relabelings.

Since there exist 24 permutations of four numbers, by choosing different permutations σ , we obtain 24 functions of the above form f_σ for a fixed function f . Some of these functions may be equal, whence the size of the conjugacy class of f may be smaller than 24. For example, if $f = \text{NOT}(a)$, then there exist only four distinct functions of the form f_σ (counting f itself). Our experiments show, however, that for the vast majority of functions, the conjugacy classes are of size 24.

2. Inversion. As mentioned above, if we know a minimal implementation for f , then we know one for its inverse as well.

Note that conjugation and inversion commute

$$(\tilde{\sigma}^{-1} \circ f \circ \tilde{\sigma})^{-1} = \tilde{\sigma}^{-1} \circ f^{-1} \circ \tilde{\sigma}.$$

For a function f , consider the union of the two conjugacy classes of f and f^{-1} . Call the elements of this union *equivalent* to f . It follows that equivalent functions have the same size. Moreover, since gates are idempotent (i.e., equal to their own inverses) and their conjugacy classes consist of gates, if we know a minimal circuit for f , we can easily obtain one for any function in the equivalence class of f . Formally, if $f = \lambda_1 \circ \dots \circ \lambda_n$, where n is the size of f and λ_i are gates, then $f^{-1} = \lambda_n \circ \dots \circ \lambda_1$, and if $f' = \tilde{\sigma}^{-1} \circ f \circ \tilde{\sigma}$, then $f' = \lambda'_1 \circ \dots \circ \lambda'_n$, where $\lambda'_i = \tilde{\sigma}^{-1} \circ \lambda_i \circ \tilde{\sigma}$ are also gates. Our experiments show that a vast majority of functions have 48 distinct equivalent functions. This fact can reduce the search space by almost a factor of 48 as follows.

For a function f , define the canonical representative of its equivalence class. A convenient canonical representative can be obtained by introducing the lexicographic order on the set of 4-bit reversible functions, considered as permutations of $\{0, 1, 2, \dots, 15\}$ and encoded accordingly by the sequence $f(0), f(1), \dots, f(15)$, and choosing the function whose corresponding sequence is lexicographically smallest. Now, instead of storing all functions of size at most k , store the canonical representative for each equivalence class. This will reduce the storage size by almost a factor of 48. Then, we use Algorithm 1 to search for a minimal circuit for a given reversible function f .

Algorithm 1. Minimal circuit (FINDOPT).

Require: Reversible function f of size at most L .

Hash table H containing canonical representatives of all equivalence classes of functions of size at most k and the last gates of their minimal circuits, $k \geq L/2$.

Lists A_i , $1 \leq i \leq L - k$, of all functions of size i .

Ensure: A minimal circuit c for f .

```

if  $f = \text{IDENTITY}$  then
    return empty circuit
end if
 $E_f \leftarrow$  equivalence class of  $f$ 
 $\bar{f} \leftarrow$  canonical representative of  $E_f$ 
if  $\bar{f} \in H$  then
     $\bar{\lambda} \leftarrow$  last gate of  $\bar{f}$ 
    if  $f$  is a conjugate of  $\bar{f}$  then
        let  $f = \tilde{\sigma}^{-1} \circ \bar{f} \circ \tilde{\sigma}$ 
         $\lambda \leftarrow \tilde{\sigma}^{-1} \circ \bar{\lambda} \circ \tilde{\sigma}$ 
         $c \leftarrow$  minimal circuit for  $f \circ \lambda$ 
        return  $c \circ \lambda$ 
    else
        let  $f = \tilde{\sigma}^{-1} \circ \bar{f}^{-1} \circ \tilde{\sigma}$ 
         $\lambda \leftarrow \tilde{\sigma}^{-1} \circ \bar{\lambda} \circ \tilde{\sigma}$ 
         $c \leftarrow$  minimal circuit for  $\lambda \circ f$ 
        return  $\lambda \circ c$ 
    end if
end if
for  $i = 1$  to  $L - k$  do
    for  $g \in A_i$  do
         $h \leftarrow g \circ f$ 
         $E_h \leftarrow$  equivalence class of  $h$ 
         $\bar{h} \leftarrow$  canonical representative of  $E_h$ 
        if  $\bar{h} \in H$  then
             $c_g \leftarrow$  minimal circuit for  $g$ 
             $c_h \leftarrow$  minimal circuit for  $h$ 
            return  $c_g^{-1} \circ c_h$ 
        end if
    end for
end for
return error: size of  $f$  is greater than  $L$ 

```

The algorithm requires a hash table with canonical representatives of equivalence classes of size at most k , together with the last gates of their minimal circuits, and lists of all permutations of size at most $L - k$. We have precomputed the canonical representatives for $k = 9$ using breadth-first search (see Algorithm 2). For efficiency reasons, we store the last *or the first* gate of a minimal

circuit for each canonical representative. However, this information is clearly sufficient to reconstruct the entire circuit and, in particular, the last gate. Using this pre-computed data, the hash table and the lists of all permutations of size at most $L - k$ are formed at the start-up. An implementation storing only the hash table is possible. Such an implementation will require less RAM memory, but it will be slower. We decided to focus on higher speed, because Table 4 indicates that we do not need to be able to search optimal circuits requiring up to $18(=9 \times 2)$ gates, which we could do, otherwise, by storing only the hash table.

Algorithm 2. Breadth-first search (BFS).

Require: k

Ensure: Lists A_i of canonical representatives of size $\leq k$; Hash table H with these canonical representatives and their first or last gates.

Let H be a hash table (keys are functions, values are gates)

$H.\text{insert}(\text{IDENTITY}, \text{HAS_NO_GATES})$

$A_0 \leftarrow \{\text{IDENTITY}\}$

for i from 1 to k **do**

for $f \in A_{i-1} \cup \{a^{-1} \mid a \in A_{i-1}\}$ **do**

for all gates λ **do**

$h \leftarrow f \circ \lambda$

$E_h \leftarrow$ equivalence class of h

$\bar{h} \leftarrow$ canonical representative of E_h

if $\bar{h} \notin H$ **then**

if h is a conjugate of \bar{h} **then**

 let $h = \tilde{\sigma}^{-1} \circ \bar{h} \circ \tilde{\sigma}$

$H.\text{insert}(\bar{h}, \tilde{\sigma}^{-1} \circ \lambda \circ \tilde{\sigma}, \text{IS_LAST_GATE})$

else

 let $h = \tilde{\sigma}^{-1} \circ \bar{h}^{-1} \circ \tilde{\sigma}$

$H.\text{insert}(\bar{h}, \tilde{\sigma}^{-1} \circ \lambda \circ \tilde{\sigma}, \text{IS_FIRST_GATE})$

end if

$A_i.\text{insert}(\bar{h})$

end if

end for

end for

end for

The correctness of Algorithm 1 is proved as follows: suppose first that the size of f is at most k . The canonical representative \bar{f} of its equivalence class will have the same size as f , so it will be found in the hash table H . Since $\bar{\lambda}$ is the last gate of a minimal circuit for \bar{f} , the size of $\bar{f} \circ \bar{\lambda}$ is one less than the size of \bar{f} . The function $f \circ \lambda$ (computed if f is a conjugate of \bar{f}) or the function $\lambda \circ f$ (computed if f is a conjugate of \bar{f}^{-1}) is equivalent to $\bar{f} \circ \bar{\lambda}$ and, therefore, also is of size one less than the size of \bar{f} . Therefore, the recursive call on that function will terminate and return a minimal circuit, which we can compose with λ (at the proper side) to obtain a minimal circuit for f . The depth of recursion is equal to the size of f , and at each call we do one hash table lookup, one computation of the canonical representative, and one conjugation of a gate (the latter can be looked up in a small table). Thus, this part of the algorithm requires negligible time.

If the size of f is greater than k , but does not exceed L , then $f = g_f \circ h$ for some h of size k and g_f of size i ,

TABLE 1
Average Times of Computing Minimal
Circuits of Sizes 0..15 (in Seconds)

Size \ k	8 (LTP)	8 (CLSTR)	9 (CLSTR)
1	8.70×10^{-7}	5.25×10^{-7}	5.23×10^{-7}
2	1.26×10^{-6}	8.32×10^{-7}	8.33×10^{-7}
3	1.66×10^{-6}	1.14×10^{-6}	1.15×10^{-6}
4	2.07×10^{-6}	1.47×10^{-6}	1.47×10^{-6}
5	2.47×10^{-6}	1.79×10^{-6}	1.79×10^{-6}
6	3.48×10^{-6}	2.11×10^{-6}	2.12×10^{-6}
7	4.22×10^{-6}	2.46×10^{-6}	2.46×10^{-6}
8	4.49×10^{-6}	2.81×10^{-6}	2.80×10^{-6}
9	1.07×10^{-5}	6.68×10^{-6}	3.11×10^{-6}
10	2.28×10^{-4}	9.31×10^{-5}	6.23×10^{-6}
11	4.27×10^{-3}	3.60×10^{-3}	7.23×10^{-5}
12	6.30×10^{-2}	5.58×10^{-2}	1.34×10^{-3}
13	4.91×10^{-1}	4.80×10^{-1}	2.20×10^{-2}
14	4.38×10^0	4.50×10^0	2.32×10^{-1}
15	N/A*	$6.14 \times 10^{1*}$	$3.61 \times 10^{0*}$

$1 \leq i \leq L - k$. Then, $g = g_f^{-1} \in A_i$. Once the inner for-loop encounters this g , it will return the minimal circuit for f , because both recursive calls are for functions of size at most k . For a function f of size $s > k$, the number of iterations required to find the minimal circuit satisfies

$$\sum_{i=1}^{s-1-k} |A_i| < r \leq \sum_{i=1}^{s-k} |A_i|.$$

At each iteration, one canonical representative is computed and looked up in the hash table. Since the size of A_i grows almost exponentially (see Table 4, left column), the search time will decrease almost exponentially, and the storage will increase exponentially, as k increases. The timings for $k = 8, 9$ measured on two different systems are summarized in Table 1 (see Section 5 for machine details). Please, note that size 15 circuits may be verified against Table 5 and consequently the time to synthesize them, for all practical purposes, is zero. We marked relevant entries in the Table 1 with an asterisk. The hash table loading and overall memory usage times were 191 seconds, 3.5 GB ($k = 8$) and 1,667 seconds, 43.04 GB ($k = 9$).

It follows from the above complexity analysis that the performance of the following key operations affect the speed most:

- composition of two functions ($f \circ g$) and inverse of a function (f^{-1}),
- computation of the canonical representative of an equivalence class, and
- hash table lookup.

In the next Section, we discuss an efficient implementation of these operations.

3.3 Implementation Details

As mentioned above, a 4-bit reversible function can be stored in a 64-bit word, by allocating 4 bits for each value of $f(0), f(1), \dots, f(15)$. Then, the composition of two functions can be computed in 94 machine instructions using the algorithm composition and the inverse function can be computed in 59 machine instructions using algorithm inverse.


```

unsigned64 composition(unsigned64 p,
                      unsigned64 q) {
    unsigned64 d = (py & 15) << 2;
    unsigned64 r = (q >> p.i) & 15;
    p >>= 2; d = p & 60;
    r |= ((q >> d) & 15) << 4;
    p >>= 4; d = p & 60;
    r |= ((q >> d) & 15) << 8;
    p >>= 4; d = p & 60;
    r |= ((q >> d) & 15) << 16;
    ...
    p >>= 4; d = p & 60;
    r |= ((q >> d) & 15) << 60;
    return r;
}

unsigned64 inverse(unsigned64 p) {
    p >>= 2;
    unsigned64 q = 1 << (p & 60);
    p >>= 4; q |= 2 << (p & 60);
    p >>= 4; q |= 3 << (p & 60);
    ...
    p >>= 4; q |= 15 << (p & 60);
    return q;
}

unsigned64 conjugate01(unsigned64 p) {
    p = (p & 0xF00FF00FF00FF00F) |
        ((p & 0x00F000F000F000F0) << 4) |
        ((p & 0x0F000F000F000F00) >> 4);
    return(p & 0xCCCCCCCCCCCCCCCC) |
        ((p & 0x1111111111111111) << 1) |
        ((p & 0x2222222222222222) >> 1);
}

```

In order to find the canonical representative in the equivalence class of a function f , we compute f^{-1} , generate all conjugates of f and f^{-1} , and choose the smallest among the resulting 48 functions. Since every permutation of $\{0, 1, 2, 3\}$ can be represented as a product of transpositions $(0, 1)$, $(1, 2)$, and $(2, 3)$, the sequence of conjugates of f by all 24 permutations can be obtained through conjugating f by these transpositions. These conjugations can be performed in 14 machine instructions each as in function `conjugate01`.

Two functions can be compared lexicographically using a single unsigned comparison of the corresponding two words. Thus, the canonical representative can be computed using one inversion, $23 \times 2 = 46$ conjugations by transpositions, and 47 comparisons, which totals to 750 machine instructions.

For the fast membership test, we use a linear probing hash table with Thomas Wang's hash function [21] (see `algorithm hash64shift`).

```

long hash64shift(long key) {
    key = (~key) + (key << 21);
    key = key ^ (key >>> 24);
    key = (key + (key << 3)) + (key << 8);
    key = key ^ (key >>> 14);
    key = (key + (key << 2)) + (key << 4);
    key = key ^ (key >>> 28);
}

```

TABLE 2
Parameters of Linear Hash Tables
Storing Canonical Representatives

	k	7	8	9
Size		2^{25}	2^{28}	2^{32}
Memory Usage		256 MB	2 GB	32 GB
Load Factor		0.58	0.84	0.51
Average Chain Length		3.14	9.18	2.63
Maximal Chain Length		92	754	86

```

key = key + (key << 31);
return key;
}

```

This function is well suited for our purposes: it is fast to compute and distributes the permutations uniformly over the hash table. The parameters of the hash tables storing the canonical representatives of equivalence classes of size k , for $k = 7, 8, 9$ are shown in Table 2.

4 SEARCHALL: AN ALGORITHM TO FIND ALL OPTIMAL CIRCUITS

We first outline our algorithm for finding all optimal circuits and then discuss it in detail in the follow up sections.

We employ a breadth first search that consists of two stages:

- Optimal circuits with 0.9 gates are found with Algorithm 2, BFS. This algorithm becomes inefficient for finding optimal circuits with 10 or more gates.
- Optimal circuits with 10 and more gates are found by storing and updating the bit vector of canonical representatives of permutations requiring a certain number of gates.

The SEARCHALL algorithm is used to find all reversible functions of size k for $k = 10, 11, \dots$, until we reach the maximal size of a reversible function. Starting from the known set of reversible functions of size 9, we consecutively proceed to sizes 10, 11, \dots . The transitions from size k to size $(k + 1)$ are carried out as follows (Sections 4.1 to 4.4).

First, we choose a compact representation for the set of reversible functions of size k , based on the following concept of an almost reduced function.

4.1 Almost Reduced Functions

Call a reversible function (permutation) p *almost reduced* if one of the following two conditions holds:

1. $p(0) = 0$ and $p(15) \in \{1, 3, 7, 15\}$.
2. $(p(0), p^{-1}(0))$ belongs to the following set

$$\{ (1, 1), (1, 2), (1, 15), (3, 1), (3, 3), (3, 4), (3, 5), (3, 12), (3, 15), (7, 1), (7, 3), (7, 7), (7, 8), (7, 9), (7, 11), (7, 15), (15, 15) \}.$$

Lemma 1. For every permutation p , there is at least one equivalent almost reduced reversible function.

Note that a reduced reversible function is not necessarily almost reduced. This will hopefully not lead to a confusion, since we are not going to deal with reduced functions in this section.

An almost reduced permutation p can be uniquely specified by the following data:

- $A_p: p(0)$.
- $B_p: p(15)$ if $p(0) = 0$, otherwise $p^{-1}(0)$.
- $Q_p: a$ permutation of 14 elements.

We call this data an *indexable specification* of the almost reduced permutation p .

The set of almost reduced reversible functions can be totally ordered by ordering their indexable specifications (A_p, B_p, Q_p) lexicographically. The *index* of an almost reduced reversible function p is defined as the number of smaller almost reversible functions. In order to compute the index of a function p , we first compute its indexable specification (A_p, B_p, Q_p) . Then, we compute $n(Q_p)$, the number of 14-element permutations smaller than Q_p , and $n(A_p, B_p)$, the number of pairs (A, B) that are a valid part of an indexable specification and are lexicographically smaller than (A_p, B_p) . Then, the index of p is given by

$$n(p) = 14!n(A_p, B_p) + n(Q_p).$$

Efficient conversions between reversible functions and their indexable specifications are quite straightforward, therefore we omit these algorithms here. Various efficient algorithms for indexing permutations are also well known.

Since almost reduced functions and their indexable specifications are in a one-to-one correspondence, the total number of almost reduced functions is $21 \cdot 14! < 1.84 \times 10^{12}$. This is ≈ 11.43 times less than the total number of reversible functions, yet about four times greater than the number of equivalence classes—i.e., different reduced permutations. The main reason why we do not index equivalence classes directly (which would have further reduced our memory requirements by about a factor of 4) is that we could not find an efficient algorithm for computing these indices.

4.2 From Size k to Size $k+1$

We encode the set of reversible functions of size k by a bit array of size $21 \times 14!$ bits (under 209 GB), where bit i is set whenever the almost reduced function p with $n(p) = i$ has size k and is the smallest almost reduced function in its equivalence class.

We further split this array in three parts called *slices*, by partitioning the set of pairs (A, B) that are valid parts of an indexable specification into three subsets. One third of the bit array easily fits in the memory of the machine we were using for the experiments (and leaves enough extra space for the system not to be tempted to turn on swapping during the computation).

Suppose that the bit array for functions of size k is stored in an input file. We compute the bit array for functions of size $k+1$ and store it in the output file via the following stages:

1. **Composition.** Repeated for each target slice s (there are three of them). Allocate in memory a bit array a of size $21 \times 14!/3$ bits. For every almost reduced function p marked in the input bit array, generate all its conjugates and inverses (thus we obtain all reversible functions of size k). Then for every

function p' in the equivalence class of p and every gate g , find an almost reduced representative q in the equivalence class of the composition $g \circ p'$, then compute its index $n(q)$. If $n(q)$ is in slice s , set the $n(q)$ th bit in the bit array a . At the end, output the array in a new file and concatenate the three slices.

2. **Canonization.** Because an equivalence class can have more than one almost reduced element, the previous stage may have marked more than one bit for some equivalence classes $g \circ p$. We scan the bit array output at the previous stage and, for each permutation q marked there, compute the smallest equivalent almost reduced permutation q' and mark the corresponding bit in the array allocated in memory. Since the entire bit array does not fit in memory, we again use three slices and at the end concatenate them.
3. **Subtraction.** The bit array produced by the previous stage contains all functions of size $k+1$ as well as some functions of size k and $k-1$. We, therefore, subtract the bit arrays corresponding to sizes k and $k-1$. The resulting bit array satisfies the property: bit i is set whenever the almost reduced function p with $n(p) = i$ has size $k+1$ and is the smallest almost reduced function in its equivalence class.
4. **(Optional Stage) Counting.** For each almost reduced function of size $k+1$, smallest in its equivalence class, we generate the entire equivalence class and count its cardinality. As a result, we obtain the total number of reversible functions of size $k+1$.

4.3 Optimization

The hardest stage to optimize is Composition. Our initial implementation, which was quite literally following the above description, with some ad hoc improvements, was going to require months to compute the functions of size 12. We found the following shortcut, which speeds it up by about a factor of 24.

For every almost reduced function p marked in the input bit array, we compute its equivalence class. However, we avoid computing the compositions of each element of the equivalence class and each gate. Instead, we extract the values $p(0), p(15), p^{-1}(0), p^{-1}(15)$. Given these values and a gate g , one can determine which conjugation and possibly inversion must be applied to $g \circ p$ to obtain an almost reduced function. The table of these conjugations and inversions is precomputed in advance.

Then, suppose a given permutation p is conjugate to an almost reduced permutation, i.e., $c^{-1} \circ g \circ p \circ c$ for some conjugation c is almost reduced. We rewrite this as $c^{-1} \circ g \circ c \circ c^{-1} \circ p \circ c$. The conjugations of the 32 gates are also precomputed in advance and stored in a separate table. Since the conjugations of p have been computed at the beginning of this step (indeed, we have computed the entire equivalence class of p), we can just take one of its elements p' and compose it with the gate $g' = c^{-1} \circ g \circ c$. The resulting permutation $g' \circ p'$ is almost reduced.

If the almost reduced representative in the equivalence class of $g \circ p$ is a conjugate of the inverse $(g \circ p)^{-1}$, i.e., equals $c^{-1} \circ (g \circ p)^{-1} \circ c$, then we rewrite this as $c^{-1} \circ p^{-1} \circ g^{-1} \circ c = c^{-1} \circ p^{-1} \circ c \circ c^{-1} \circ g \circ c$ (also using the fact that

$g^{-1} = g$ for every gate g). Now, we again observe that $c^{-1} \circ p^{-1} \circ c$ has been precomputed, so we only need to compose it with the gate $g' = c^{-1} \circ g \circ c$. Note that compositions of functions with gates (on either side) can be performed very efficiently.

Having implemented this optimization, we were able to compute all reversible functions of size 10. However, the computation of functions of size 12 would still take too long, so we parallelized the algorithm.

4.4 Parallelization

Both composition and canonization stages are computationally intensive. We parallelized them using the following architecture implemented with MPI.

For composition, the master job reads the input bit vector in blocks. Every block is sent to one of 16 workers, which are chosen in a circular (round robin) order. These workers decode the bits in the blocks into permutations of size k , apply gates to them as described above, and compute the indices of the resulting almost reduced permutations. These indices are stored in a temporary array, which is partitioned into eight equal slices. Once all indices have been computed by a worker, the slices are sent to the corresponding eight collector jobs. Each collector possesses its own bit vector allocated in RAM. It receives arrays with indices of bits to be marked from the 16 workers in a round robin order. Having received an array from a worker, it marks the corresponding bits. At the end of a round, the collector signals the master that a round has been completed. At the very end, the collectors write their bit vectors to disk in sequence.

The master makes sure that the collectors are not more than 80 blocks behind it. If it continued to send the blocks to the workers without waiting for the workers and collectors to finish processing them, the unprocessed blocks would accumulate in the communication channels between the master and the workers. This results in a memory leak, which turned out to be faster than the system swapping mechanism, and therefore, caused a deadlock. By allowing the collectors to be only a certain number of blocks behind the master, we restrict the amount of data in the communication channels at any given moment and, thus, prevent the leak. It is useful to allow a nonzero lag though, for otherwise the system becomes overly synchronized, which drastically reduces the performance: the workers and collectors that finish first end up waiting on the others most of the time. With the lag, communication channels work as buffers, from which the workers continue to draw data. The amount of data in each particular channel at a given moment may vary, depending on the speed at which the corresponding worker processes the previous blocks.

Exactly the same parallel architecture is used for canonization. The master again reads the input bit vector in blocks. The workers compute the minimal almost reduced equivalent permutation for each almost reduced permutation they receive from the master and send their indices to the collectors. The collectors mark the corresponding bits and write those bit vectors to disk at the end.

TABLE 3
Distribution of the Number of Gates Required for
10,000,000 Random 4-Bit Reversible Functions

Size	Functions
14	17,191
13	2,371,039
12	5,110,943
11	2,051,507
10	392,108
9	50,861
8	5,269
7	455
6	24
5	3

5 PERFORMANCE AND RESULTS

We performed several tests using two computer systems, LPTP and CLSTR. LPTP is a Sony VGN-NS190D laptop with Intel Core Duo 2.0 GHz processor, 4 GB RAM, and a 5400 RPM SATA HDD running Linux. CLSTR is a cluster [22] located at the Institute for Quantum Computing. We used a single Sun X4600 node with 128 GB RAM and 8 AMD Opteron quad-core CPUs for each run of the SEARCHALL and FINDOPT algorithm in CLSTR. The following sections summarize the tests and results.

5.1 Synthesis of Random Permutations

In this test, we generated 10,000,000 random uniformly distributed permutations using the Mersenne twister random number generator [10]. We next generated their optimal circuits using algorithm FINDOPT. The test was executed on CLSTR. It took 75,613.12 seconds (about 21 hours) of user time and the maximal RAM memory usage was 43.04 GB. Note that 1,667 seconds (approximately 28 minutes) were spent loading previously computed optimal circuits with up to nine gates (see Section 5.2 for details) into RAM. On average, it took only 0.00756 seconds to synthesize an optimal circuit for a permutation. The distribution of the circuit sizes is shown in Table 3.

Since there are no permutations requiring 16 or more gates, and only a few permutations requiring 15 gates (see Section 5.3 for details), this implies that the search FINDOPT may be easily modified to explicitly store all optimal 15-bit implementations in the cache, and search optimal implementations with up to 14 gates. Such a search may be executed using a computer capable of storing reduced optimal implementations with up to seven gates, i.e., a machine with only 256MB of available RAM. In other words, FINDOPT allows performing optimal 4-bit circuit calculation even on an older machine.

5.2 Distribution of Optimal Implementations

Table 4 lists the distribution of the number of permutations that can be realized with optimal circuits requiring a specified number of gates. We used CLSTR to run this test, and it took 1,130,276 seconds (approximately 13 days) to complete it. Circuits with up to nine gates were synthesized using the BFS algorithm. For circuits with 10 gates and more we used SEARCHALL.

TABLE 4
Number of 4-Bit Permutations
Requiring Prescribed Number of Gates

Size	Functions	Reduced Functions	Runtime
≥ 16	0	0	
15	144	5	66,782s
14	37,481,795,636	781,068,573	245,488s
13	4,959,760,623,552	103,331,100,613	397,464s
12	10,690,104,057,901	222,714,352,278	238,589s
11	4,298,462,792,398	89,554,073,333	103,595s
10	819,182,578,179	17,067,688,249	68,670s
9	105,984,823,653	2,208,511,226	8,836.36s
8	10,804,681,959	225,242,556	744.41s
7	932,651,938	19,466,575	95.574s
6	70,763,560	1,482,686	11.109s
5	4,807,552	101,983	0.816s
4	294,507	6,538	0.06s
3	16,204	425	0.004s
2	784	33	<0.001s
1	32	4	<0.001s
0	1	1	<0.001s
Total	20,922,789,888,000	435,903,095,078	1,130,276s

We have calculated the average number of gates required for a random 4-bit reversible function, 11.93937....

5.3 Most Complex Permutations

As follows from the previous section, there are only five reduced permutations requiring the maximal number of gates, 15. We list all five canonical representatives, together with their optimal implementations, in Table 5. Columns of this table report the function specification, the number of symmetries this specification generates, and an optimal circuit found by our program, correspondingly. The remaining 139 ($= 144 - 5$) permutations requiring 15 gates in an optimal implementation may be found via reducing them to a canonical representative through an input/output relabeling and possible inversion. For example, [6, 8, 15, 13, 4, 0, 12, 1, 3, 9, 11, 14, 10, 2, 5, 7] is a permutation requiring 15 gates in an optimal implementation. It may be obtained from the third listed in the Table 5 via inversion and relabeling $(a, b, c, d) \mapsto (c, a, b, d)$.

5.4 Optimal Linear Circuits

Linear reversible circuits are the most complex part of quantum error correcting circuits [1]. Efficiency of these circuits defines the efficiency of quantum encoding and decoding error correction operations. Linear reversible

TABLE 6
Number of 4-Bit Linear Reversible Functions
Requiring 0..10 Gates in an Optimal Implementation

Size	Functions
10	138
9	13,555
8	84,225
7	118,424
6	72,062
5	26,182
4	6,589
3	1,206
2	162
1	16
0	1

functions are those whose positive polarity Reed-Muller polynomial has only linear terms. More simply, and equivalently, linear reversible functions are those computable by circuits with NOT and CNOT gates.

For example, the reversible mapping $a, b, c, d \mapsto b \oplus 1, a \oplus c \oplus 1, d \oplus 1, a$ is a linear reversible function. Interestingly, this linear function is one of the 138 most complex linear reversible functions—it requires 10 gates in an optimal implementation. The optimal implementation of this function is given by the circuit CNOT(b,a) CNOT(c,d) CNOT(d,b) NOT(d) CNOT(a,b) CNOT(d,c) CNOT(b,d) CNOT(d,a) NOT(d) CNOT(c,b).

We synthesized optimal circuits for all 322,560 4-bit linear reversible functions using the FINDOPT algorithm. This process took under 2 seconds on LPTP. The distribution of the number of functions requiring a given number of gates is shown in Table 6.

5.5 Synthesis of Benchmarks

In this section, we report optimal circuits for benchmark functions that have been previously reported in the literature. Table 7 summarizes the results. The table describes the **Name** of the benchmark function, its complete **Specification**, **Size** of the **Best Known Circuit (SBKC)**, the **Source** of this circuit, indicator of whether this circuit has been **Proved Optimal (PO?)**, **Size** of an **Optimal Circuit (SOC)**, the optimal implementation that our program found, and the runtime our program takes to find this optimal implementation. We used the head node of CLSTR for this test, and report the runtime it takes after hash table with all optimal implementations with up to nine gates is loaded into RAM. Shorter runtimes were identified using

TABLE 5
Permutations Requiring 15 Gates

Function	# Symm.	Implementation
[1,5,0,8,9,11,2,15,3,12,4,6,10,14,13,7]	24	CNOT(a,c) CNOT(c,d) CNOT(d,a) TOF(b,d,c) CNOT(a,b) TOF(c,d,b) TOF4(a,b,c,d) CNOT(c,a) NOT(b) NOT(c) CNOT(a,d) TOF(b,d,c) TOF(b,c,a) TOF(a,c,b) NOT(c)
[1,9,0,4,10,8,2,11,3,15,5,12,7,14,13,6]	24	NOT(d) CNOT(d,c) TOF4(a,c,d,b) TOF(a,d,c) TOF(b,d,a) TOF(c,d,b) TOF(b,c,d) TOF(a,d,b) CNOT(a,d) NOT(a) NOT(b) NOT(c) TOF4(b,c,d,a) CNOT(b,c) TOF(a,d,c)
[3,1,7,13,11,0,8,15,2,5,10,6,9,14,12,4]	48	NOT(b) CNOT(b,a) TOF(a,b,c) TOF(a,d,b) CNOT(c,d) TOF4(b,c,d,a) TOF4(a,b,c,d) CNOT(a,c) CNOT(c,b) TOF(b,d,c) NOT(a) NOT(b) CNOT(c,d) CNOT(d,a) TOF(a,b,c)
[3,1,11,7,8,0,9,5,2,6,15,13,14,4,10,12]	24	CNOT(c,b) CNOT(a,d) CNOT(d,a) TOF4(a,b,c,d) TOF(a,b,c) TOF(b,c,a) TOF(a,d,b) CNOT(b,c) NOT(d) NOT(c) NOT(a) TOF(c,d,b) TOF(b,c,d) CNOT(d,c) CNOT(a,c)
[3,5,11,1,8,0,9,7,2,6,14,13,10,4,12,15]	24	CNOT(c,b) TOF(b,d,a) CNOT(a,d) CNOT(d,c) TOF(b,c,a) TOF(a,c,b) TOF(a,d,c) TOF(b,c,a) NOT(d) NOT(c) NOT(b) CNOT(d,a) TOF(b,c,d) CNOT(d,b) TOF(a,b,c)

TABLE 7
Optimal Implementations of Benchmark Functions

Name	Specification	SBKC	Source	PO?	SOC	Our optimal circuit	Runtime
4_49	[15,1,12,3,5,6,8,7,0,10,13,9,2,4,14,11]	12	[7]	No	12	NOT(a) CNOT(c,a) CNOT(a,d) TOF(a,b,d) CNOT(d,a) TOF(c,d,b) TOF(a,d,c) TOF(b,c,a) TOF(a,b,d) NOT(a) CNOT(d,b) CNOT(d,c)	.000355s
4bit-7-8	[0,1,2,3,4,5,6,8,7,9,10,11,12,13,14,15]	7	[11]	No	7	CNOT(d,b) CNOT(d,a) CNOT(c,d) TOF4(a,b,d,c) CNOT(c,d) CNOT(d,b) CNOT(d,a)	.000002s
decode42	[1,2,4,8,0,3,5,6,7,9,10,11,12,13,14,15]	11	[4]	No	10	CNOT(c,b) CNOT(d,a) CNOT(c,a) TOF(a,d,b) CNOT(b,c) TOF4(a,b,c,d) TOF(b,d,c) CNOT(c,a) CNOT(a,b) NOT(a)	.000004s
hwb4	[0,2,4,12,8,5,9,11,1,6,10,13,3,14,7,15]	11	[7]	Yes	11	CNOT(b,d) CNOT(d,a) CNOT(a,c) CNOT(c,d) TOF(a,d,b) TOF(b,c,a) CNOT(d,c) CNOT(c,b) TOF(a,c,b) CNOT(a,c) CNOT(b,d)	.000052s
imark	[4,5,2,14,0,3,6,10,11,8,15,1,12,13,7,9]	7	[16]	No	7	TOF(c,d,a) TOF(a,b,d) CNOT(d,c) CNOT(b,c) CNOT(d,a) TOF(a,c,b) NOT(c)	.000003s
mperk	[3,11,2,10,0,7,1,6,15,8,14,9,13,5,12,4]	9*	[11], [15]	No	9	NOT(c) CNOT(d,c) TOF(c,d,b) TOF(a,c,d) CNOT(b,a) CNOT(d,a) CNOT(c,a) CNOT(a,b) CNOT(b,c)	.000003s
oc5	[6,0,12,15,7,1,5,2,4,10,13,3,11,8,14,9]	15	[17]	No	11	TOF(b,d,c) TOF(c,d,b) TOF(a,b,c) NOT(a) CNOT(d,b) CNOT(c,a) CNOT(d,c) TOF(a,b,d) CNOT(c,a) CNOT(c,b) TOF4(a,b,d,c)	.000158s
oc6	[9,0,2,15,11,6,7,8,14,3,4,13,5,1,12,10]	14	[17]	No	12	TOF4(a,b,c,d) TOF(b,d,c) CNOT(d,a) TOF(b,c,d) CNOT(c,b) CNOT(b,c) TOF(a,d,c) TOF(b,c,a) TOF(a,b,c) NOT(a) CNOT(d,b) CNOT(a,d)	.000380s
oc7	[6,15,9,5,13,12,3,7,2,10,1,11,0,14,4,8]	17	[17]	No	13	CNOT(b,d) NOT(b) TOF(a,b,c) TOF(b,d,a) TOF(c,d,b) CNOT(a,d) CNOT(a,c) CNOT(b,a) TOF4(a,b,c,d) TOF(c,d,b) CNOT(c,a) NOT(a) CNOT(b,c)	.0194s
oc8	[11,3,9,2,7,13,15,14,8,1,4,10,0,12,6,5]	16	[17]	No	12	CNOT(a,b) TOF(b,c,a) TOF(c,d,b) CNOT(d,a) TOF4(a,b,d,c) TOF(a,b,d) NOT(b) TOF(a,d,b) TOF(b,d,a) TOF(b,c,d) NOT(a) CNOT(a,d)	.000725s
nth_prime4_inc	[0,2,3,5,7,11,13,1,4,6,8,9,10,12,14,15]	N/A	N/A	N/A	11	TOF(a,b,c) CNOT(d,b) TOF(a,c,b) TOF(b,d,c) TOF(b,c,d) CNOT(a,b) TOF4(b,c,d,a) CNOT(c,b) TOF4(a,b,d,c) CNOT(b,a) TOF(b,d,a)	0.000095s
rd32	[0,7,6,9,4,11,10,13,8,15,14,1,12,3,2,5]	4	[2]	Yes	4	TOF(a,b,d) CNOT(a,b) TOF(b,c,d) CNOT(b,c)	.000001s
shift4	[1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,0]	4	[11]	Yes	4	TOF4(a,b,c,d) TOF(a,b,c) CNOT(a,b) NOT(a)	.000002s

multiple runs of the search to achieve sufficient accuracy. Please note that we introduce the function *nth_prime4_inc*, which cannot be found in the previous literature. Also, the nine-gate circuit for the function *mperk* reported in [15] requires some extra SWAP gates to properly map inputs into their respective outputs, indicated by an asterisk.

6 CONCLUSIONS AND POSSIBLE EXTENSIONS

In this paper, we described two algorithms: first, FINDOPT, finds an optimal circuit for any 4-bit reversible function, and second, SEARCALL, finds all optimal 4-bit reversible circuits. Our goal was to minimize the number of gates required for function implementation. Our implementation of FINDOPT takes approximately 3 hours to calculate all optimal implementations requiring up to nine gates, and then an average of about 0.00756 seconds to search for an optimal circuit of any 4-bit reversible function. Our implementation of SEARCALL requires about about 13 days, however, it needs to be completed only once to collect all relevant statistics and data. Both calculations are surprisingly fast given the size of the search space.

Using BFS, we demonstrated the synthesis of 117,798,040,190 optimal circuits in 9,688 seconds, amounting to an average speed of 12,168,356 circuits per second. This is over 70 times faster and some 4,500 times more than the best previously reported result (26 million circuits in 152 seconds) [16]. Furthermore, using FINDALL, we demonstrated the

synthesis of 20,922,789,888,000 functions in 1,130,276 seconds (18,511,222 circuits per second). This is over 100 times faster and over 800,000 times more than in [16].

We also demonstrated that the search for any given optimal circuit can be done very quickly—0.00756 seconds per a random function. For example, if all optimal circuits were written into a *hypothetical* 100+TB 5400 RPM hard drive, the average time to extract a random circuit from the drive would be expected to take on the order of 0.01-0.02 seconds (typical access time for 5400 RPM hard drives). In other words, it would take longer to read the answer from a *hypothetical* hard drive than to compute it with our implementation. Furthermore, the 3-hour calculation of all optimal circuits with up to nine gates could be reduced by storing its result (computed once for the entirety of the described search and its follow up executions) on the hard drive, as was done in Section 5.1. It took 1,667 seconds, i.e., under 28 minutes, to load optimal circuits with up to nine gates into RAM using CLSTR. Given that the media transfer rate of modern hard drives is 1 Gbit/s (=1 GB in 8 seconds) and higher, it may take no longer than 5 minutes (=300 s > 296 s = 37 * 8 s) to load optimal implementations into RAM to initiate the search on a different machine.

Minor modifications to the algorithm could be explored to address other optimization issues. For example, for practicality, one may be interested in minimizing depth. This may be important if a faster circuit is preferred, and/or if quantum noise has a stronger constituent with time,

TABLE 8
Distribution of the Number of Gates Required for
LNN-Optimal Implementation of 10,000,000
Random 4-Bit Reversible Functions

Size	Functions
18	6
17	20,546
16	1,091,953
15	3,976,746
14	3,286,497
13	1,244,670
12	308,993
11	59,289
10	9,693
9	1,387
8	189
7	26
6	5

than with the disturbance introduced by multiple gate applications. It may also be important to account for the different implementation costs of the gates used (generally, NOT is much simpler than CNOT, which in turn, is simpler than Toffoli). Both modifications are possible, by making minor changes to the first part of FINDOPT, and minor modifications of SEARCHALL. To optimize depth, one needs to consider a different family of gates, where, for instance, sequence $\text{NOT}(a) \text{CNOT}(b, c)$ is counted as a single gate. To account for different gate costs, one needs to search for small circuits via increasing cost by one (assuming costs are given as natural numbers), as opposed to adding a gate to all maximal size optimal circuits.

It is also possible to extend the search to find optimal implementations in restricted architectures (see the Section 7 for details). Finally, the search could be extended to find some small optimal 5-bit circuits. A simple calculation shows that $80^6 \log_2(80)/5!/2$ bits (the number of elementary transformations to the power of depth, times space to store a single gate, divided by the number of symmetries) suffices to store all optimal circuits containing up to six gates for 5-bit permutations. Thus, a search of optimal implementations may be carried to compute optimal circuits with up to 12 gates. However, it is possible that a larger search may be performed.

Finally, techniques reported in this paper may be applied to the synthesis of optimal stabilizer circuits. Coupled with peep-hole optimization algorithm for circuit simplification, these results may become a very useful tool in optimizing error correction circuits. This may be of a particular practical interest since implementations of quantum algorithms may be expected to be dominated by the error correction circuits.

7 EXTENSION: LNN CIRCUITS

Of all possible extensions of the presented search algorithms described in Section 6, the most computationally difficult is the one where the underlying architecture is restricted. This is because the number of input/output labeling symmetries that can be used to reduce the search is equal to the number automorphisms of the unlabeled graph corresponding to the underlying architecture. For the complete graph on four bits, K_4 , its number of automorphisms is maximal, $4! = 24$. Since the number of automorphisms is maximized, this has helped us to gain maximal advantage. Of the connected

TABLE 9
Number of 4-Bit Permutations Requiring Prescribed
Number of Gates in the LNN Architecture

Size	Functions	Reduced Functions	Runtime
≥ 11	Unknown	Unknown	
10	20,355,134,386	5,089,090,158	13,299.9s
9	2,921,376,642	730,451,187	1,642.72s
8	378,041,753	94,551,844	241.367s
7	44,754,539	11,201,218	68.192s
6	4,886,991	1,226,080	11.041s
5	493,788	124,628	1.28s
4	46,108	11,885	0.18s
3	3,947	1,083	0.02s
2	303	100	<0.001s
1	20	10	<0.001s
0	1	1	<0.001s
Total	23,704,738,478	5,926,658,194	15,264.692s

graphs with four nodes, the chain, corresponding to the linear nearest neighbor (LNN) architecture, has the least number of automorphisms, being just two. In this section, we will illustrate that our search may be modified to find optimal circuits in the LNN architecture, implying that it is at least as efficient for the remaining four possible architectures (the number of nonisomorphic unlabeled connected graphs on four nodes minus two, one for the LNN and one for K_4).

The restriction to the LNN architecture implies that the gate set is limited to those gates operating on the subset of qubits that is a continuous substring of the string of all variables, $abcd$. For example, gates $\text{CNOT}(a, b)$ and $\text{TOF}(b, d, c)$ are allowed, and gate $\text{TOF}(a, b, d)$ is disallowed. Such restriction to the LNN architecture does not necessarily imply direct physical applications. In fact, not only it is not certain that the underlying architecture is LNN (and, it must be noted that local architecture may differ from global architecture), and not only do we not account for the individual gate costs, but the restriction itself may not be physically grounded. Indeed, according to [9], LNN-optimal NCV implementation of the $\text{TOF}(a, c, b)$ requires 13 gates and depth 13; however, $\text{TOF}(a, b, c)$ requires only nine gates and depth 9, and $\text{TOF}(a, b, d)$ may be implemented with 15 NCV gates and depth of only 12. This means that $\text{TOF}(a, b, d)$ is “faster” than $\text{TOF}(a, c, b)$, and of the two we have just disallowed $\text{TOF}(a, b, d)$. In other words, we suggest that our software is updated to achieve the results relevant to experiments once all physical restrictions are known; however, the goal of this paper, and this Section, in particular, is to illustrate that when needed proper modifications are possible.

We have modified implementation of the FINDOPT algorithm to account for restrictions imposed by the LNN architecture. This required to change the definition of the equivalence class of a 4-bit reversible function. In particular, the newly defined LNN-equivalence class allows symmetries with respect to the inversion, and one of the two possible relabelings: $(a, b, c, d) \mapsto (a, b, c, d)$ and $(a, b, c, d) \mapsto (d, c, b, a)$. The remainder of code and algorithms remained essentially the same.

We report the result of the optimal LNN synthesis in the following three tables, Tables 8, 9, and 10. The data contained is analogous to that reported in Tables 3, 4, and 7 correspondingly. It took the total of 1,313,020.23 seconds to

TABLE 10
LNN-Optimal Implementations of Benchmark Functions

Name	# gates	LNN optimal circuit	Runtime
4_49	16	CNOT(d,c) TOF(c,d,b) CNOT(b,a) NOT(c) TOF(b,c,d) CNOT(c,b) TOF4(a,b,c,d) TOF4(a,c,d,b) TOF(a,b,c) CNOT(a,b) NOT(a) CNOT(c,d) TOF(b,d,c) TOF(b,c,a) TOF(a,c,b) NOT(c)	0.68s
4bit-7-8	7	CNOT(d,c) TOF(c,d,b) TOF4(b,c,d,a) TOF4(a,b,c,d) TOF4(b,c,d,a) TOF(c,d,b) CNOT(d,c)	<0.001s
decode42	13	NOT(d) NOT(c) CNOT(c,b) TOF(b,d,c) TOF(b,c,a) TOF4(a,c,d,b) TOF4(a,b,c,d) CNOT(d,c) TOF(b,c,a) TOF4(a,c,d,b) NOT(d) TOF(c,d,b) NOT(c)	<0.001s
hwb4	16	TOF(a,b,c) CNOT(c,d) CNOT(b,c) TOF(b,c,a) TOF4(a,c,d,b) NOT(a) CNOT(a,b) CNOT(d,c) CNOT(b,a) TOF(b,c,d) TOF(b,c,a) TOF4(a,c,d,b) TOF(a,b,c) NOT(b) CNOT(c,d) TOF(a,b,c)	1.34s
imark	11	TOF4(b,c,d,a) TOF4(a,b,c,d) CNOT(b,c) TOF4(a,b,c,d) TOF4(b,c,d,a) CNOT(d,c) TOF(a,c,b) CNOT(b,a) TOF(c,d,b) NOT(c) CNOT(b,a)	<0.001s
mperk	11	CNOT(d,c) TOF(a,c,b) CNOT(b,a) TOF(c,d,b) CNOT(a,b) NOT(c) TOF(b,c,d) CNOT(c,b) CNOT(b,a) TOF(c,d,b) CNOT(b,c)	<0.001s
oc5	14	CNOT(d,c) TOF4(a,c,d,b) CNOT(c,b) TOF(b,d,c) CNOT(b,a) CNOT(b,c) TOF(b,c,d) CNOT(c,b) NOT(a) CNOT(a,b) CNOT(b,c) TOF(b,c,a) TOF4(b,c,d,a) TOF4(a,b,d,c)	0.01s
oc6	14	TOF(b,c,a) TOF4(a,c,d,b) TOF4(a,b,c,d) NOT(b) CNOT(b,a) CNOT(d,c) TOF4(b,c,d,a) CNOT(c,b) TOF(a,b,c) CNOT(c,d) TOF(b,d,c) NOT(b) TOF(a,b,c) TOF(b,c,d)	0.01s
oc7	15	TOF(b,c,d) TOF(c,d,b) TOF4(b,c,d,a) TOF(a,b,c) NOT(b) TOF4(a,b,d,c) TOF(b,c,d) CNOT(c,b) CNOT(b,a) NOT(d) TOF(c,d,b) TOF4(a,b,d,c) NOT(a) CNOT(c,d) TOF(a,b,c)	0.07s
oc8	14	CNOT(a,b) TOF(b,d,c) CNOT(c,d) NOT(c) TOF(b,c,a) TOF4(a,b,c,d) CNOT(a,b) CNOT(d,c) CNOT(b,c) TOF4(a,c,d,b) CNOT(c,b) CNOT(b,a) TOF4(a,b,c,d) TOF(b,d,c)	0.03s
nth_prime4	11	CNOT(d,c) TOF(b,c,a) CNOT(b,c) NOT(b) TOF(b,c,d) TOF(b,c,a) TOF4(a,b,d,c) TOF(a,c,b) NOT(a) TOF4(a,c,d,b) CNOT(b,a)	<0.001s
rd32	7	TOF(b,c,d) NOT(c) TOF4(a,b,c,d) NOT(c) CNOT(a,b) TOF4(a,b,c,d) CNOT(b,c)	<0.001s
shift4	4	TOF4(a,b,c,d) TOF(a,b,c) CNOT(a,b) NOT(a)	<0.001s

calculate 10,000,000 LNN-optimal reversible circuits reported in the Table 8. This calculation has been performed on three cluster nodes in parallel, reducing physical time spent on this calculation by a factor of 3. The average time to calculate a single random LNN-optimal 4-bit reversible circuit is approximately 0.131 seconds. Based on the number and distribution of circuits in Tables 4 and 9 we conjecture that there are no LNN-optimal circuits requiring 21 gates, and as such it suffices to have generated all optimal circuits with up to 10 gates to synthesize any LNN-optimal 4-bit reversible circuit.

ACKNOWLEDGMENTS

The authors wish to thank Dr. Donny Cheung from Tornado Medical Systems and Dr. Sean Falconer from Stanford University for their useful discussions and contributions. We wish to thank Marek Szyprowski and Professor Pawel Kerntopf from Warsaw University of Technology for their comments. This article was based on work partially supported by the US National Science Foundation (NSF), during D. Maslov's assignment at the Foundation.

REFERENCES

- [1] S. Aaronson and D. Gottesman, "Improved Simulation of Stabilizer Circuits," *Physical Rev. A*, vol. 70, no. 5, p. 052328, <http://arxiv.org/abs/quant-ph/0406196>, 2004.
- [2] R.P. Feynman, "Quantum Mechanical Computers," *Foundations of Physics*, vol. 16, no. 6, pp. 507-531, June 1986.
- [3] D. Große et al., "Exact SAT-Based Toffoli Network Synthesis," *Proc. 17th ACM Great Lakes Symp. VLSI*, 2007.
- [4] P. Gupta et al., "An Algorithm for Synthesis of Reversible Logic Circuits," *IEEE Trans. Computer-Aided Design of Integrated Circuits and Systems*, vol. 25, no. 11, pp. 2317-2330, Nov. 2006.
- [5] H. Häffner et al., "Scalable Multiparticle Entanglement of Trapped Ions," *Nature*, vol. 438, pp. 643-646, <http://arxiv.org/abs/quant-ph/0603217>, Dec. 2005.
- [6] P. Kerntopf, "A New Heuristic Algorithm for Reversible Logic Synthesis," *Proc. IEEE Design Automation Conf. (DAC)*, pp. 834-837, 2004.
- [7] D. Maslov, "Reversible Logic Synthesis Benchmarks Page," <http://webhome.cs.uvic.ca/~dmaslov/>, Jan. 2011.
- [8] D. Maslov et al., "Techniques for the Synthesis of Reversible Toffoli Networks," *ACM Trans. Design Automation of Electronic Systems*, vol. 12, no. 4, pp. 1-28, <http://arxiv.org/abs/quant-ph/0607166>, Sept. 2007.
- [9] D. Maslov and D.M. Miller, "Comparison of the Cost Metrics for Reversible and Quantum Logic Synthesis," *IET Computers and Digital Techniques*, vol. 1, no. 2, pp. 98-104, <http://arxiv.org/abs/quant-ph/0511008>, 2007.
- [10] M. Matsumoto and T. Nishimura, "Mersenne Twister: A 623-Dimensionally Equidistributed Uniform Pseudo-Random Number Generator," *ACM Trans. Modeling and Computer Simulation*, vol. 8, no. 1, pp. 3-30, 1998.
- [11] D.M. Miller, "Spectral and Two-Place Decomposition Techniques in Reversible Logic," *Proc. 45th Midwest Symp. Circuits and Systems (MWSCAS)*, Aug. 2002.
- [12] M. Nielsen and I. Chuang, *Quantum Computation and Quantum Information*. Cambridge Univ. Press, 2000.
- [13] C. Negrevergne et al., "Benchmarking Quantum Control Methods on a 12-Qubit System," *Physical Rev. Letters*, vol. 96, no. 17, p. 170501, <http://arxiv.org/abs/quant-ph/0603248>, 2006.
- [14] A. Peres, "Reversible Logic and Quantum Computers," *Physical Rev. A*, vol. 32, no. 6, pp. 3266-3276, 1985.
- [15] M. Perkowski et al., "A General Decomposition for Reversible Logic," *Proc. Fifth Reed-Muller Workshop*, pp. 119-138, 2001.
- [16] A.K. Prasad et al., "Algorithms and Data Structures for Simplifying Reversible Circuits," *ACM J. Emerging Technologies in Computing Systems*, vol. 2, no. 4, pp. 277-293, 2006.
- [17] M. Saeedi et al., "On the Behavior of Substitution-Based Reversible Circuit Synthesis Algorithms Investigation and Improvement," *Proc. IEEE Computer Soc. Ann. Symp. VLSI (ISVLSI '07)*, pp. 428-436, Mar. 2007.
- [18] V.V. Shende et al., "Synthesis of Reversible Logic Circuits," *IEEE Trans. Computer-Aided Design of Integrated Circuits and Systems*, vol. 22, no. 6, pp. 710-722, <http://arxiv.org/abs/quant-ph/0207001>, June 2003.
- [19] P.W. Shor, "Polynomial-Time Algorithms for Prime Factorization and Discrete Logarithms on a Quantum Computer," *SIAM J. Computing*, vol. 26, pp. 1484-1509, 1997.
- [20] G. Yang et al., "Bi-Directional Synthesis of 4-Bit Reversible Circuits," *Computer J.*, vol. 51, no. 2, pp. 207-215, Mar. 2008.
- [21] T. Wang, "Integer Hash Function, Version 3.1," <http://www.concentric.net/~Ttwang/tech/inthash.htm>, Mar. 2007.
- [22] IQC Cluster, <http://www.iqc.ca/computing/cluster.html>, Sept. 2010.



Oleg Golubitsky received the PhD degree in mathematics from Lomonosov's Moscow State University, Russia and the PhD degree in computer science from the University of New Brunswick, Canada. Currently, he is working as a software engineer at Google, Inc., and has worked as an assistant professor of computer science at the University of New Brunswick and as a postdoctoral fellow at the University of Pisa, Italy, Queen's University, Canada, and the

University of Western Ontario, Canada. His publications areas include computer algebra, differential algebra, handwriting recognition, and quantum computing.



Dmitri Maslov received the MSc degree in mathematics from Lomonosov's Moscow State University, Russia, in 2000, and the MSc and PhD degrees in computer science from the University of New Brunswick, Canada, in 2002 and 2003, respectively. Currently, he is working as a program director in the Division of Computing and Communication Foundations, directorate for Computer and Information Science and Engineering, US National Science

Foundation, and as an adjunct faculty in the Department of Physics and Astronomy, the University of Waterloo, and the Institute for Quantum Computing, Canada. His current research interests include quantum circuits and architectures, quantum algorithms, quantum information processing, logic synthesis, reversible logic, and circuit complexity. He is a member of the IEEE.

▷ **For more information on this or any other computing topic, please visit our Digital Library at www.computer.org/publications/dlib.**